

June 25, 2023

1 Wpływ transfer learningu na sieci neuronowe.

Autorzy: Adam Klekowski, Paweł Kruczkiewicz, Piotr Makarewicz

Transfer learning to technika, w której model uczenia maszynowego wykorzystuje wiedzę i umiejętności nabyte na jednym zadaniu do rozwiązania innego, pokrewnego zadania. Wykonuje się ją poprzez dodanie warstwy gęsto połączonej na koniec modelu.

Fine-tuning to proces dostosowywania modelu nauczonego przy użyciu transfer learningu poprzez dotrenowanie kilku ostatnich bądź wszystkich warstw do konkretnej domeny lub problemu. Robi się to w celu uzyskania lepszej wydajności dla zbioru innego niż oryginalny.

Te techniki stosuje się, aby przyspieszyć proces uczenia, wykorzystać wiedzę z dużych zbiorów danych lub zamodelować złożone relacje, gdy nie ma wystarczająco dużo danych do treningu modelu od podstaw.

1.1 Przygotowanie

1.1.1 Import bibliotek

Do zrealizowania zadania użyliśmy biblioteki tensorflow. Obliczenia wykonywano na

```
[ ]: import os
import time
import shutil
import random

import numpy as np
import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: gpus = tf.config.list_physical_devices('GPU')
print("Num GPUs Available: ", len(gpus))
if len(gpus) > 0:
    tf.config.experimental.set_memory_growth(gpus[0], True)
```

Num GPUs Available: 1

1.1.2 Funkcje pomocnicze

```
[ ]: ## consts
DATA_FOLDER="data-trimmed"

INPUT_SHAPE = (256, 256, 3)
BATCH_SIZE = 32
```

```
[ ]: class TimeHistory(keras.callbacks.Callback):
    def on_train_begin(self, logs={}):
        self.times = []

    def on_epoch_begin(self, epoch, logs={}):
        self.epoch_time_start = time.time()

    def on_epoch_end(self, epoch, logs={}):
        self.times.append(time.time() - self.epoch_time_start)

def plot_train_results(history, times, title):
    h = history.history
    loss, acc, val_loss, val_acc = h["loss"], h["categorical_accuracy"],
    ↪h["val_loss"], h["val_categorical_accuracy"]
    avg_epoch_time = np.round(np.mean(times), 1)
    x = np.arange(len(loss)) + 1

    fig, (ax1, ax2) = plt.subplots(1, 2)
    fig.suptitle(f"{title} [Avg epoch time: {avg_epoch_time} s]")

    ax1.set_title("accuracy")
    ax1.plot(x, acc, label="Test")
    ax1.plot(x, val_acc, label="Validation")
    ax1.set_ylim([0, 1])
    ax1.legend()

    ax2.set_title("loss")
    ax2.plot(x, loss, label="Test")
    ax2.plot(x, val_loss, label="Validation")
    ax2.set_ylim([0, 5])
    ax2.legend()

def train_and_check_model(model, model_name, t_ds, v_ds, epochs=150,
    ↪lr_scale=1):
    """Trains the model and plots training results"""
    time_history = TimeHistory()

    model.compile(optimizer=keras.optimizers.Nadam(learning_rate=0.
    ↪001*lr_scale),
```

```

        loss="categorical_crossentropy", metrics=[keras.metrics.
↪CategoricalAccuracy()])
    history = model.fit(t_ds,
                        validation_data=v_ds,
                        batch_size=BATCH_SIZE,
                        epochs=epochs,
                        callbacks=[time_history])

    model.save(f"./models/{model_name}")
    plot_train_results(history, time_history.times, model_name)

```

1.1.3 Pobranie i przygotowanie zbioru

Wybraliśmy zbiór Food-101, który zawiera zdjęcia 101 różnych rodzajów jedzenia. Z tego zbioru wybraliśmy 20 klas i z każdej klasy wzięliśmy 256 próbek. Następnie zbiór podzielono na zbiór treningowy i walidacyjny w proporcji 4:1.

```

[ ]: ! wget http://data.vision.ee.ethz.ch/cvl/food-101.tar.gz
     ! tar xzvf food-101.tar.gz
     ! mv food-101 data & rm food-101.tar.gz

```

```

[ ]: SAMPLES_PER_CLASS = 256
     OLD_FOLDER = "data"
     NEW_FOLDER = "data-trimmed"
     CHOSEN_CLASSES = {'apple_pie', 'pizza', 'hamburger', 'spaghetti_bolognese',
↪ 'chocolate_cake', 'hot_dog', 'ice_cream',
                        'carrot_cake', 'chicken_curry', 'churros', 'falafel',
↪ 'fish_and_chips', 'french_fries', 'hummus',
                        'greek_salad', 'panna_cotta', 'nachos', 'lasagna', 'tacos',
↪ 'risotto'}
     NUM_CLASSES = len(CHOSEN_CLASSES)

     print(f'There are {NUM_CLASSES} chosen classes:')
     for idx, class_name in enumerate(CHOSEN_CLASSES):
         print(f"{idx:2}. {class_name}")

     if not os.path.exists(NEW_FOLDER):
         os.mkdir(NEW_FOLDER)

     for class_label in CHOSEN_CLASSES:
         old_class_dir = f"{OLD_FOLDER}/{class_label}"
         new_class_dir = f"{NEW_FOLDER}/{class_label}"

         if not os.path.exists(new_class_dir):
             os.mkdir(new_class_dir)
             # print(f"Creating {new_class_dir}")

```

```

list_of_samples = os.listdir(old_class_dir)
trimmed_samples = random.sample(list_of_samples, SAMPLES_PER_CLASS)
# print(f"Number of samples: {len(trimmed_samples)}")

for sample_name in trimmed_samples:
    shutil.copyfile(f"{old_class_dir}/{sample_name}", f"{new_class_dir}/
↪{sample_name}")
    # print(f"Copying {sample_name} to {new_class_dir}")

```

There are 20 chosen classes:

0. lasagna
1. french_fries
2. hamburger
3. carrot_cake
4. hot_dog
5. panna_cotta
6. greek_salad
7. chicken_curry
8. hummus
9. apple_pie
10. fish_and_chips
11. nachos
12. ice_cream
13. spaghetti_bolognese
14. churros
15. chocolate_cake
16. pizza
17. tacos
18. falafel
19. risotto

```

[ ]: train_ds, validation_ds = keras.utils.image_dataset_from_directory(
    directory=DATA_FOLDER,
    label_mode='categorical',
    image_size=(256, 256),
    validation_split=0.2,
    subset="both",
    seed=21)

```

Found 5120 files belonging to 20 classes.

Using 4096 files for training.

Using 1024 files for validation.

```

[ ]: labels = sorted(os.listdir(DATA_FOLDER))

plt.figure(figsize=(10,10))

for (batch_of_images, batch_of_labels) in train_ds.take(1):

```

```

batch_of_images = batch_of_images[:9]
batch_of_labels = batch_of_labels[:9]
for i, (image, inferred_label) in enumerate(zip(batch_of_images,
↪batch_of_labels)):
    image = image / 255
    label = labels[np.nonzero(np.array(inferred_label))[0][0]]
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(image)
    plt.title(label)
    plt.axis("off")

```

chicken_curry



fish_and_chips



carrot_cake



nachos



hamburger



greek_salad



carrot_cake



risotto



fish_and_chips



1.2 EfficientNetV2B0

1.2.1 Użycie EfficientNetV2B0

Jako model wykorzystaliśmy model EfficientNet wytrenowany na zbiorze ImageNet.

```
[ ]: efficient_net = keras.applications.EfficientNetV2B0(  
    weights='imagenet', # Load weights pre-trained on ImageNet.  
    input_shape=(256, 256, 3),  
    include_top=False,  
    include_preprocessing=False) # Do not include the ImageNet classifier at  
    ↪ the top.  
  
efficient_net.summary()
```

Model: "efficientnetv2-b0"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 256, 256, 3 0)]		[]
stem_conv (Conv2D) ['input_2[0][0]']	(None, 128, 128, 32 864)		
stem_bn (BatchNormalization) ['stem_conv[0][0]']	(None, 128, 128, 32 128)		
stem_activation (Activation) ['stem_bn[0][0]']	(None, 128, 128, 32 0)		
block1a_project_conv (Conv2D) ['stem_activation[0][0]']	(None, 128, 128, 16 4608)		
input_2 (InputLayer)	[(None, 256, 256, 3 0)]		[]
stem_conv (Conv2D) ['input_2[0][0]']	(None, 128, 128, 32 864)		

```

    stem_bn (BatchNormalization) (None, 128, 128, 32 128
['stem_conv[0][0]']
    )

    stem_activation (Activation) (None, 128, 128, 32 0
['stem_bn[0][0]']
    )

    block1a_project_conv (Conv2D) (None, 128, 128, 16 4608
['stem_activation[0][0]']
    )

    block1a_project_bn (BatchNorma (None, 128, 128, 16 64
['block1a_project_conv[0][0]']
    lization)
    )

    block1a_project_activation (Ac (None, 128, 128, 16 0
['block1a_project_bn[0][0]']
    tivation)
    )

    block2a_expand_conv (Conv2D) (None, 64, 64, 64) 9216
['block1a_project_activation[0][0]
    ]']

    block2a_expand_bn (BatchNormal (None, 64, 64, 64) 256
['block2a_expand_conv[0][0]']
    ization)

    block2a_expand_activation (Act (None, 64, 64, 64) 0
['block2a_expand_bn[0][0]']
    ivation)

    block2a_project_conv (Conv2D) (None, 64, 64, 32) 2048
['block2a_expand_activation[0][0]
    ']'

    block2a_project_bn (BatchNorma (None, 64, 64, 32) 128
['block2a_project_conv[0][0]']
    lization)

    block2b_expand_conv (Conv2D) (None, 64, 64, 128) 36864
['block2a_project_bn[0][0]']

    block2b_expand_bn (BatchNormal (None, 64, 64, 128) 512
['block2b_expand_conv[0][0]']
    ization)

```

```

    block2b_expand_activation (Act (None, 64, 64, 128) 0
['block2b_expand_bn[0][0]']
    ivation)

    block2b_project_conv (Conv2D) (None, 64, 64, 32) 4096
['block2b_expand_activation[0][0]

    ]

    block2b_project_bn (BatchNorma (None, 64, 64, 32) 128
['block2b_project_conv[0][0]']
    lization)

    block2b_drop (Dropout) (None, 64, 64, 32) 0
['block2b_project_bn[0][0]']

    block2b_add (Add) (None, 64, 64, 32) 0
['block2b_drop[0][0]',
'block2a_project_bn[0][0]']

    block3a_expand_conv (Conv2D) (None, 32, 32, 128) 36864
['block2b_add[0][0]']

    block3a_expand_bn (BatchNormal (None, 32, 32, 128) 512
['block3a_expand_conv[0][0]']
    ization)

    block3a_expand_activation (Act (None, 32, 32, 128) 0
['block3a_expand_bn[0][0]']
    ivation)

    block3a_project_conv (Conv2D) (None, 32, 32, 48) 6144
['block3a_expand_activation[0][0]

    ]

    block3a_project_bn (BatchNorma (None, 32, 32, 48) 192
['block3a_project_conv[0][0]']
    lization)

    block3b_expand_conv (Conv2D) (None, 32, 32, 192) 82944
['block3a_project_bn[0][0]']

    block3b_expand_bn (BatchNormal (None, 32, 32, 192) 768
['block3b_expand_conv[0][0]']
    ization)

    block3b_expand_activation (Act (None, 32, 32, 192) 0
['block3b_expand_bn[0][0]']
    ivation)

```



```

    block3b_project_conv (Conv2D) (None, 32, 32, 48) 9216
['block3b_expand_activation[0][0]

    block3b_project_bn (BatchNormal (None, 32, 32, 48) 192
['block3b_project_conv[0][0] 'l
    ization)

    block3b_drop (Dropout) (None, 32, 32, 48) 0
['block3b_project_bn[0][0] 'l

    block3b_add (Add) (None, 32, 32, 48) 0
['block3b_drop[0][0] ',
'block3a_project_bn[0][0] 'l

    block4a_expand_conv (Conv2D) (None, 32, 32, 192) 9216
['block3b_add[0][0] 'l

    block4a_expand_bn (BatchNormal (None, 32, 32, 192) 768
['block4a_expand_conv[0][0] 'l
    ization)

    block4a_expand_activation (Act (None, 32, 32, 192) 0
['block4a_expand_bn[0][0] 'l
    ivation)

    block4a_dwconv2 (DepthwiseConv (None, 16, 16, 192) 1728
['block4a_expand_activation[0][0]
    2D)

    block4a_bn (BatchNormalization (None, 16, 16, 192) 768
['block4a_dwconv2[0][0] 'l
    )

    block4a_activation (Activation (None, 16, 16, 192) 0
['block4a_bn[0][0] 'l
    )

    block4a_se_squeeze (GlobalAver (None, 192) 0
['block4a_activation[0][0] 'l
    agePooling2D)

    block4a_se_reshape (Reshape) (None, 1, 1, 192) 0
['block4a_se_squeeze[0][0] 'l

    block4a_se_reduce (Conv2D) (None, 1, 1, 12) 2316
['block4a_se_reshape[0][0] 'l

```

block4a_se_expand (Conv2D)	(None, 1, 1, 192)	2496
['block4a_se_reduce[0][0]']		
block4a_se_excite (Multiply)	(None, 16, 16, 192)	0
['block4a_activation[0][0]', 'block4a_se_expand[0][0]']		
block4a_project_conv (Conv2D)	(None, 16, 16, 96)	18432
['block4a_se_excite[0][0]']		
block4a_project_bn (BatchNormal	(None, 16, 16, 96)	384
['block4a_project_conv[0][0]'] lization)		
block4b_expand_conv (Conv2D)	(None, 16, 16, 384)	36864
['block4a_project_bn[0][0]']		
block4b_expand_bn (BatchNormal	(None, 16, 16, 384)	1536
['block4b_expand_conv[0][0]'] ization)		
block4b_expand_activation (Act	(None, 16, 16, 384)	0
['block4b_expand_bn[0][0]'] ivation)		
block4b_dwconv2 (DepthwiseConv	(None, 16, 16, 384)	3456
['block4b_expand_activation[0][0]'] 2D)		
block4b_bn (BatchNormalization	(None, 16, 16, 384)	1536
['block4b_dwconv2[0][0]'])		
block4b_activation (Activation	(None, 16, 16, 384)	0
['block4b_bn[0][0]'])		
block4b_se_squeeze (GlobalAver	(None, 384)	0
['block4b_activation[0][0]'] agePooling2D)		
block4b_se_reshape (Reshape)	(None, 1, 1, 384)	0
['block4b_se_squeeze[0][0]']		
block4b_se_reduce (Conv2D)	(None, 1, 1, 24)	9240
['block4b_se_reshape[0][0]']		

```

    block4b_se_expand (Conv2D)      (None, 1, 1, 384)    9600
['block4b_se_reduce[0][0]']

    block4b_se_excite (Multiply)    (None, 16, 16, 384)  0
['block4b_activation[0][0]',
'block4b_se_expand[0][0]']

    block4b_project_conv (Conv2D)   (None, 16, 16, 96)   36864
['block4b_se_excite[0][0]']

    block4b_project_bn (BatchNorma  (None, 16, 16, 96)   384
['block4b_project_conv[0][0]'
lization)

    block4b_drop (Dropout)          (None, 16, 16, 96)   0
['block4b_project_bn[0][0]']

    block4b_add (Add)               (None, 16, 16, 96)   0
['block4b_drop[0][0]',
'block4a_project_bn[0][0]']

    block4c_expand_conv (Conv2D)    (None, 16, 16, 384)  36864
['block4b_add[0][0]']

    block4c_expand_bn (BatchNormal  (None, 16, 16, 384)  1536
['block4c_expand_conv[0][0]'
lization)

    block4c_expand_activation (Act  (None, 16, 16, 384)  0
['block4c_expand_bn[0][0]'
ivation)

    block4c_dwconv2 (DepthwiseConv  (None, 16, 16, 384)  3456
['block4c_expand_activation[0][0]
2D)

    block4c_bn (BatchNormalization  (None, 16, 16, 384)  1536
['block4c_dwconv2[0][0]']
)

    block4c_activation (Activation  (None, 16, 16, 384)  0
['block4c_bn[0][0]']
)

    block4c_se_squeeze (GlobalAver  (None, 384)          0
['block4c_activation[0][0]'
agePooling2D)

```

block4c_se_reshape (Reshape)	(None, 1, 1, 384)	0
['block4c_se_squeeze[0][0]']		
block4c_se_reduce (Conv2D)	(None, 1, 1, 24)	9240
['block4c_se_reshape[0][0]']		
block4c_se_expand (Conv2D)	(None, 1, 1, 384)	9600
['block4c_se_reduce[0][0]']		
block4c_se_excite (Multiply)	(None, 16, 16, 384)	0
['block4c_activation[0][0]', 'block4c_se_expand[0][0]']		
block4c_project_conv (Conv2D)	(None, 16, 16, 96)	36864
['block4c_se_excite[0][0]']		
block4c_project_bn (BatchNormal	(None, 16, 16, 96)	384
['block4c_project_conv[0][0]'] lization)		
block4c_drop (Dropout)	(None, 16, 16, 96)	0
['block4c_project_bn[0][0]']		
block4c_add (Add)	(None, 16, 16, 96)	0
['block4c_drop[0][0]', 'block4b_add[0][0]']		
block5a_expand_conv (Conv2D)	(None, 16, 16, 576)	55296
['block4c_add[0][0]']		
block5a_expand_bn (BatchNormal	(None, 16, 16, 576)	2304
['block5a_expand_conv[0][0]'] lization)		
block5a_expand_activation (Act	(None, 16, 16, 576)	0
['block5a_expand_bn[0][0]'] ivation)		
block5a_dwconv2 (DepthwiseConv	(None, 16, 16, 576)	5184
['block5a_expand_activation[0][0]'] 2D)		
block5a_bn (BatchNormalization	(None, 16, 16, 576)	2304
['block5a_dwconv2[0][0]'])		
block5a_activation (Activation	(None, 16, 16, 576)	0
['block5a_bn[0][0]']		

```

)

block5a_se_squeeze (GlobalAveragePooling2D) (None, 576) 0
['block5a_activation[0][0]']

block5a_se_reshape (Reshape) (None, 1, 1, 576) 0
['block5a_se_squeeze[0][0]']

block5a_se_reduce (Conv2D) (None, 1, 1, 24) 13848
['block5a_se_reshape[0][0]']

block5a_se_expand (Conv2D) (None, 1, 1, 576) 14400
['block5a_se_reduce[0][0]']

block5a_se_excite (Multiply) (None, 16, 16, 576) 0
['block5a_activation[0][0]',
'block5a_se_expand[0][0]']

block5a_project_conv (Conv2D) (None, 16, 16, 112) 64512
['block5a_se_excite[0][0]']

block5a_project_bn (BatchNormalization) (None, 16, 16, 112) 448
['block5a_project_conv[0][0]']

block5b_expand_conv (Conv2D) (None, 16, 16, 672) 75264
['block5a_project_bn[0][0]']

block5b_expand_bn (BatchNormalization) (None, 16, 16, 672) 2688
['block5b_expand_conv[0][0]']

block5b_expand_activation (Activation) (None, 16, 16, 672) 0
['block5b_expand_bn[0][0]']

block5b_dwconv2 (DepthwiseConv2D) (None, 16, 16, 672) 6048
['block5b_expand_activation[0][0]']

block5b_bn (BatchNormalization) (None, 16, 16, 672) 2688
['block5b_dwconv2[0][0]']

block5b_activation (Activation) (None, 16, 16, 672) 0
['block5b_bn[0][0]']

```

block5b_se_squeeze (GlobalAveragePooling2D)	(None, 672)	0
['block5b_activation[0][0]']		
block5b_se_reshape (Reshape)	(None, 1, 1, 672)	0
['block5b_se_squeeze[0][0]']		
block5b_se_reduce (Conv2D)	(None, 1, 1, 28)	18844
['block5b_se_reshape[0][0]']		
block5b_se_expand (Conv2D)	(None, 1, 1, 672)	19488
['block5b_se_reduce[0][0]']		
block5b_se_excite (Multiply)	(None, 16, 16, 672)	0
['block5b_activation[0][0]', 'block5b_se_expand[0][0]']		
block5b_project_conv (Conv2D)	(None, 16, 16, 112)	75264
['block5b_se_excite[0][0]']		
block5b_project_bn (BatchNormalization)	(None, 16, 16, 112)	448
['block5b_project_conv[0][0]']		
block5b_drop (Dropout)	(None, 16, 16, 112)	0
['block5b_project_bn[0][0]']		
block5b_add (Add)	(None, 16, 16, 112)	0
['block5b_drop[0][0]', 'block5a_project_bn[0][0]']		
block5c_expand_conv (Conv2D)	(None, 16, 16, 672)	75264
['block5b_add[0][0]']		
block5c_expand_bn (BatchNormalization)	(None, 16, 16, 672)	2688
['block5c_expand_conv[0][0]']		
block5c_expand_activation (Activation)	(None, 16, 16, 672)	0
['block5c_expand_bn[0][0]']		
block5c_dwconv2 (DepthwiseConv2D)	(None, 16, 16, 672)	6048
['block5c_expand_activation[0][0]']		
block5c_bn (BatchNormalization)	(None, 16, 16, 672)	2688

```

['block5c_dwconv2[0][0]']
)

block5c_activation (Activation (None, 16, 16, 672) 0
['block5c_bn[0][0]']
)

block5c_se_squeeze (GlobalAveragePooling2D) (None, 672) 0
['block5c_activation[0][0]']

block5c_se_reshape (Reshape) (None, 1, 1, 672) 0
['block5c_se_squeeze[0][0]']

block5c_se_reduce (Conv2D) (None, 1, 1, 28) 18844
['block5c_se_reshape[0][0]']

block5c_se_expand (Conv2D) (None, 1, 1, 672) 19488
['block5c_se_reduce[0][0]']

block5c_se_excite (Multiply) (None, 16, 16, 672) 0
['block5c_activation[0][0]',
'block5c_se_expand[0][0]']

block5c_project_conv (Conv2D) (None, 16, 16, 112) 75264
['block5c_se_excite[0][0]']

block5c_project_bn (BatchNormalization) (None, 16, 16, 112) 448
['block5c_project_conv[0][0]']

block5c_drop (Dropout) (None, 16, 16, 112) 0
['block5c_project_bn[0][0]']

block5c_add (Add) (None, 16, 16, 112) 0
['block5c_drop[0][0]',
'block5b_add[0][0]']

block5d_expand_conv (Conv2D) (None, 16, 16, 672) 75264
['block5c_add[0][0]']

block5d_expand_bn (BatchNormalization) (None, 16, 16, 672) 2688
['block5d_expand_conv[0][0]']

block5d_expand_activation (Activation) (None, 16, 16, 672) 0
['block5d_expand_bn[0][0]']

```

```

block5d_dwconv2 (DepthwiseConv (None, 16, 16, 672) 6048
['block5d_expand_activation[0][0]
2D)

block5d_bn (BatchNormalization (None, 16, 16, 672) 2688
['block5d_dwconv2[0][0]')

block5d_activation (Activation (None, 16, 16, 672) 0
['block5d_bn[0][0]')

block5d_se_squeeze (GlobalAveragePooling2D) (None, 672) 0
['block5d_activation[0][0]']

block5d_se_reshape (Reshape) (None, 1, 1, 672) 0
['block5d_se_squeeze[0][0]']

block5d_se_reduce (Conv2D) (None, 1, 1, 28) 18844
['block5d_se_reshape[0][0]']

block5d_se_expand (Conv2D) (None, 1, 1, 672) 19488
['block5d_se_reduce[0][0]']

block5d_se_excite (Multiply) (None, 16, 16, 672) 0
['block5d_activation[0][0]',
'block5d_se_expand[0][0]']

block5d_project_conv (Conv2D) (None, 16, 16, 112) 75264
['block5d_se_excite[0][0]']

block5d_project_bn (BatchNormalization) (None, 16, 16, 112) 448
['block5d_project_conv[0][0]']

block5d_drop (Dropout) (None, 16, 16, 112) 0
['block5d_project_bn[0][0]']

block5d_add (Add) (None, 16, 16, 112) 0
['block5d_drop[0][0]',
'block5c_add[0][0]']

block5e_expand_conv (Conv2D) (None, 16, 16, 672) 75264
['block5d_add[0][0]']

block5e_expand_bn (BatchNormalization) (None, 16, 16, 672) 2688

```



```

['block5e_expand_conv[0][0]']
ization)

block5e_expand_activation (Activation (None, 16, 16, 672) 0
['block5e_expand_bn[0][0]']
ivation)

block5e_dwconv2 (DepthwiseConv (None, 16, 16, 672) 6048
['block5e_expand_activation[0][0]
2D)

block5e_bn (BatchNormalization (None, 16, 16, 672) 2688
['block5e_dwconv2[0][0]']
)

block5e_activation (Activation (None, 16, 16, 672) 0
['block5e_bn[0][0]']
)

block5e_se_squeeze (GlobalAveragePooling2D (None, 672) 0
['block5e_activation[0][0]']
agePooling2D)

block5e_se_reshape (Reshape (None, 1, 1, 672) 0
['block5e_se_squeeze[0][0]']

block5e_se_reduce (Conv2D (None, 1, 1, 28) 18844
['block5e_se_reshape[0][0]']

block5e_se_expand (Conv2D (None, 1, 1, 672) 19488
['block5e_se_reduce[0][0]']

block5e_se_excite (Multiply (None, 16, 16, 672) 0
['block5e_activation[0][0]'],
['block5e_se_expand[0][0]']

block5e_project_conv (Conv2D (None, 16, 16, 112) 75264
['block5e_se_excite[0][0]']

block5e_project_bn (BatchNormalization (None, 16, 16, 112) 448
['block5e_project_conv[0][0]']
lization)

block5e_drop (Dropout (None, 16, 16, 112) 0
['block5e_project_bn[0][0]']

block5e_add (Add (None, 16, 16, 112) 0
['block5e_drop[0][0]'],

```

```

'block5d_add[0][0]']

block6a_expand_conv (Conv2D)    (None, 16, 16, 672)  75264
['block5e_add[0][0]']

block6a_expand_bn (BatchNormal  (None, 16, 16, 672)  2688
['block6a_expand_conv[0][0]']
ization)

block6a_expand_activation (Act  (None, 16, 16, 672)  0
['block6a_expand_bn[0][0]']
ivation)

block6a_dwconv2 (DepthwiseConv  (None, 8, 8, 672)  6048
['block6a_expand_activation[0][0]
2D)

block6a_bn (BatchNormalization  (None, 8, 8, 672)  2688
['block6a_dwconv2[0][0]']
)

block6a_activation (Activation  (None, 8, 8, 672)  0
['block6a_bn[0][0]']
)

block6a_se_squeeze (GlobalAver  (None, 672)  0
['block6a_activation[0][0]']
agePooling2D)

block6a_se_reshape (Reshape)    (None, 1, 1, 672)  0
['block6a_se_squeeze[0][0]']

block6a_se_reduce (Conv2D)      (None, 1, 1, 28)  18844
['block6a_se_reshape[0][0]']

block6a_se_expand (Conv2D)      (None, 1, 1, 672)  19488
['block6a_se_reduce[0][0]']

block6a_se_excite (Multiply)    (None, 8, 8, 672)  0
['block6a_activation[0][0]',
'block6a_se_expand[0][0]']

block6a_project_conv (Conv2D)   (None, 8, 8, 192)  129024
['block6a_se_excite[0][0]']

block6a_project_bn (BatchNorma  (None, 8, 8, 192)  768
['block6a_project_conv[0][0]']
lization)

```

block6b_expand_conv (Conv2D)	(None, 8, 8, 1152)	221184	
['block6a_project_bn[0][0]']			
block6b_expand_bn (BatchNormal	(None, 8, 8, 1152)	4608	
['block6b_expand_conv[0][0]']			
ization)			
block6b_expand_activation (Act	(None, 8, 8, 1152)	0	
['block6b_expand_bn[0][0]']			
ivation)			
block6b_dwconv2 (DepthwiseConv	(None, 8, 8, 1152)	10368	
['block6b_expand_activation[0][0]			
2D)			']
block6b_bn (BatchNormalization	(None, 8, 8, 1152)	4608	
['block6b_dwconv2[0][0]']			
)			
block6b_activation (Activation	(None, 8, 8, 1152)	0	
['block6b_bn[0][0]']			
)			
block6b_se_squeeze (GlobalAver	(None, 1152)	0	
['block6b_activation[0][0]']			
agePooling2D)			
block6b_se_reshape (Reshape)	(None, 1, 1, 1152)	0	
['block6b_se_squeeze[0][0]']			
block6b_se_reduce (Conv2D)	(None, 1, 1, 48)	55344	
['block6b_se_reshape[0][0]']			
block6b_se_expand (Conv2D)	(None, 1, 1, 1152)	56448	
['block6b_se_reduce[0][0]']			
block6b_se_excite (Multiply)	(None, 8, 8, 1152)	0	
['block6b_activation[0][0]',			
'block6b_se_expand[0][0]']			
block6b_project_conv (Conv2D)	(None, 8, 8, 192)	221184	
['block6b_se_excite[0][0]']			
block6b_project_bn (BatchNorma	(None, 8, 8, 192)	768	
['block6b_project_conv[0][0]']			
lization)			

block6b_drop (Dropout)	(None, 8, 8, 192)	0
['block6b_project_bn[0][0]']		
block6b_add (Add)	(None, 8, 8, 192)	0
['block6b_drop[0][0]',		
'block6a_project_bn[0][0]']		
block6c_expand_conv (Conv2D)	(None, 8, 8, 1152)	221184
['block6b_add[0][0]']		
block6c_expand_bn (BatchNormal	(None, 8, 8, 1152)	4608
ization)		
block6c_expand_activation (Act	(None, 8, 8, 1152)	0
ivation)		
block6c_dwconv2 (DepthwiseConv	(None, 8, 8, 1152)	10368
2D)		
block6c_bn (BatchNormalization	(None, 8, 8, 1152)	4608
)		
block6c_activation (Activation	(None, 8, 8, 1152)	0
)		
block6c_se_squeeze (GlobalAver	(None, 1152)	0
agePooling2D)		
block6c_se_reshape (Reshape)	(None, 1, 1, 1152)	0
block6c_se_reduce (Conv2D)	(None, 1, 1, 48)	55344
block6c_se_expand (Conv2D)	(None, 1, 1, 1152)	56448
block6c_se_excite (Multiply)	(None, 8, 8, 1152)	0
block6c_project_conv (Conv2D)	(None, 8, 8, 192)	221184

```

['block6c_se_excite[0][0]']

block6c_project_bn (BatchNormal (None, 8, 8, 192) 768
['block6c_project_conv[0][0]']
lization)

block6c_drop (Dropout) (None, 8, 8, 192) 0
['block6c_project_bn[0][0]']

block6c_add (Add) (None, 8, 8, 192) 0
['block6c_drop[0][0]',
'block6b_add[0][0]']

block6d_expand_conv (Conv2D) (None, 8, 8, 1152) 221184
['block6c_add[0][0]']

block6d_expand_bn (BatchNormal (None, 8, 8, 1152) 4608
['block6d_expand_conv[0][0]']
ization)

block6d_expand_activation (Act (None, 8, 8, 1152) 0
['block6d_expand_bn[0][0]']
ivation)

block6d_dwconv2 (DepthwiseConv (None, 8, 8, 1152) 10368
['block6d_expand_activation[0][0]
2D)

block6d_bn (BatchNormalization (None, 8, 8, 1152) 4608
['block6d_dwconv2[0][0]']
)

block6d_activation (Activation (None, 8, 8, 1152) 0
['block6d_bn[0][0]']
)

block6d_se_squeeze (GlobalAver (None, 1152) 0
['block6d_activation[0][0]']
agePooling2D)

block6d_se_reshape (Reshape) (None, 1, 1, 1152) 0
['block6d_se_squeeze[0][0]']

block6d_se_reduce (Conv2D) (None, 1, 1, 48) 55344
['block6d_se_reshape[0][0]']

block6d_se_expand (Conv2D) (None, 1, 1, 1152) 56448
['block6d_se_reduce[0][0]']

```

block6d_se_excite (Multiply)	(None, 8, 8, 1152)	0
['block6d_activation[0][0]', 'block6d_se_expand[0][0]']		
block6d_project_conv (Conv2D)	(None, 8, 8, 192)	221184
['block6d_se_excite[0][0]']		
block6d_project_bn (BatchNormal	(None, 8, 8, 192)	768
['block6d_project_conv[0][0]'] lization)		
block6d_drop (Dropout)	(None, 8, 8, 192)	0
['block6d_project_bn[0][0]']		
block6d_add (Add)	(None, 8, 8, 192)	0
['block6d_drop[0][0]', 'block6c_add[0][0]']		
block6e_expand_conv (Conv2D)	(None, 8, 8, 1152)	221184
['block6d_add[0][0]']		
block6e_expand_bn (BatchNormal	(None, 8, 8, 1152)	4608
['block6e_expand_conv[0][0]'] ization)		
block6e_expand_activation (Act	(None, 8, 8, 1152)	0
['block6e_expand_bn[0][0]'] ivation)		
block6e_dwconv2 (DepthwiseConv	(None, 8, 8, 1152)	10368
['block6e_expand_activation[0][0]' 2D)		
block6e_bn (BatchNormalization	(None, 8, 8, 1152)	4608
['block6e_dwconv2[0][0]'])		
block6e_activation (Activation	(None, 8, 8, 1152)	0
['block6e_bn[0][0]'])		
block6e_se_squeeze (GlobalAver	(None, 1152)	0
['block6e_activation[0][0]'] agePooling2D)		
block6e_se_reshape (Reshape)	(None, 1, 1, 1152)	0
['block6e_se_squeeze[0][0]']		

block6e_se_reduce (Conv2D) ['block6e_se_reshape[0][0]']	(None, 1, 1, 48)	55344
block6e_se_expand (Conv2D) ['block6e_se_reduce[0][0]']	(None, 1, 1, 1152)	56448
block6e_se_excite (Multiply) ['block6e_activation[0][0]', 'block6e_se_expand[0][0]']	(None, 8, 8, 1152)	0
block6e_project_conv (Conv2D) ['block6e_se_excite[0][0]']	(None, 8, 8, 192)	221184
block6e_project_bn (BatchNormal lization) ['block6e_project_conv[0][0]']	(None, 8, 8, 192)	768
block6e_drop (Dropout) ['block6e_project_bn[0][0]']	(None, 8, 8, 192)	0
block6e_add (Add) ['block6e_drop[0][0]', 'block6d_add[0][0]']	(None, 8, 8, 192)	0
block6f_expand_conv (Conv2D) ['block6e_add[0][0]']	(None, 8, 8, 1152)	221184
block6f_expand_bn (BatchNormal ization) ['block6f_expand_conv[0][0]']	(None, 8, 8, 1152)	4608
block6f_expand_activation (Act ivation) ['block6f_expand_bn[0][0]']	(None, 8, 8, 1152)	0
block6f_dwconv2 (DepthwiseConv 2D) ['block6f_expand_activation[0][0]']	(None, 8, 8, 1152)	10368
block6f_bn (BatchNormalization) ['block6f_dwconv2[0][0]']	(None, 8, 8, 1152)	4608
block6f_activation (Activation) ['block6f_bn[0][0]']	(None, 8, 8, 1152)	0

block6f_se_squeeze (GlobalAveragePooling2D)	(None, 1152)	0	
['block6f_activation[0][0]']			
block6f_se_reshape (Reshape)	(None, 1, 1, 1152)	0	
['block6f_se_squeeze[0][0]']			
block6f_se_reduce (Conv2D)	(None, 1, 1, 48)	55344	
['block6f_se_reshape[0][0]']			
block6f_se_expand (Conv2D)	(None, 1, 1, 1152)	56448	
['block6f_se_reduce[0][0]']			
block6f_se_excite (Multiply)	(None, 8, 8, 1152)	0	
['block6f_activation[0][0]', 'block6f_se_expand[0][0]']			
block6f_project_conv (Conv2D)	(None, 8, 8, 192)	221184	
['block6f_se_excite[0][0]']			
block6f_project_bn (BatchNormalization)	(None, 8, 8, 192)	768	
['block6f_project_conv[0][0]']			
block6f_drop (Dropout)	(None, 8, 8, 192)	0	
['block6f_project_bn[0][0]']			
block6f_add (Add)	(None, 8, 8, 192)	0	
['block6f_drop[0][0]', 'block6e_add[0][0]']			
block6g_expand_conv (Conv2D)	(None, 8, 8, 1152)	221184	
['block6f_add[0][0]']			
block6g_expand_bn (BatchNormalization)	(None, 8, 8, 1152)	4608	
['block6g_expand_conv[0][0]']			
block6g_expand_activation (Activation)	(None, 8, 8, 1152)	0	
['block6g_expand_bn[0][0]']			
block6g_dwconv2 (DepthwiseConv2D)	(None, 8, 8, 1152)	10368	
['block6g_expand_activation[0][0]']			
block6g_bn (BatchNormalization)	(None, 8, 8, 1152)	4608	
['block6g_dwconv2[0][0]']			


```

)

block6g_activation (Activation (None, 8, 8, 1152) 0
['block6g_bn[0][0]']
)

block6g_se_squeeze (GlobalAveragePooling2D) (None, 1152) 0
['block6g_activation[0][0]']

block6g_se_reshape (Reshape) (None, 1, 1, 1152) 0
['block6g_se_squeeze[0][0]']

block6g_se_reduce (Conv2D) (None, 1, 1, 48) 55344
['block6g_se_reshape[0][0]']

block6g_se_expand (Conv2D) (None, 1, 1, 1152) 56448
['block6g_se_reduce[0][0]']

block6g_se_excite (Multiply) (None, 8, 8, 1152) 0
['block6g_activation[0][0]',
'block6g_se_expand[0][0]']

block6g_project_conv (Conv2D) (None, 8, 8, 192) 221184
['block6g_se_excite[0][0]']

block6g_project_bn (BatchNormalization) (None, 8, 8, 192) 768
['block6g_project_conv[0][0]']

block6g_drop (Dropout) (None, 8, 8, 192) 0
['block6g_project_bn[0][0]']

block6g_add (Add) (None, 8, 8, 192) 0
['block6g_drop[0][0]',
'block6f_add[0][0]']

block6h_expand_conv (Conv2D) (None, 8, 8, 1152) 221184
['block6g_add[0][0]']

block6h_expand_bn (BatchNormalization) (None, 8, 8, 1152) 4608
['block6h_expand_conv[0][0]']

block6h_expand_activation (Activation) (None, 8, 8, 1152) 0
['block6h_expand_bn[0][0]']

```

```

    block6h_dwconv2 (DepthwiseConv (None, 8, 8, 1152) 10368
['block6h_expand_activation[0][0]
2D)

    block6h_bn (BatchNormalization (None, 8, 8, 1152) 4608
['block6h_dwconv2[0][0]')

    block6h_activation (Activation (None, 8, 8, 1152) 0
['block6h_bn[0][0]')

    block6h_se_squeeze (GlobalAveragePooling2D) (None, 1152) 0
['block6h_activation[0][0]']

    block6h_se_reshape (Reshape) (None, 1, 1, 1152) 0
['block6h_se_squeeze[0][0]']

    block6h_se_reduce (Conv2D) (None, 1, 1, 48) 55344
['block6h_se_reshape[0][0]']

    block6h_se_expand (Conv2D) (None, 1, 1, 1152) 56448
['block6h_se_reduce[0][0]']

    block6h_se_excite (Multiply) (None, 8, 8, 1152) 0
['block6h_activation[0][0]',
'block6h_se_expand[0][0]']

    block6h_project_conv (Conv2D) (None, 8, 8, 192) 221184
['block6h_se_excite[0][0]']

    block6h_project_bn (BatchNormalization) (None, 8, 8, 192) 768
['block6h_project_conv[0][0]']

    block6h_drop (Dropout) (None, 8, 8, 192) 0
['block6h_project_bn[0][0]']

    block6h_add (Add) (None, 8, 8, 192) 0
['block6h_drop[0][0]',
'block6g_add[0][0]']

    top_conv (Conv2D) (None, 8, 8, 1280) 245760
['block6h_add[0][0]']

    top_bn (BatchNormalization) (None, 8, 8, 1280) 5120
['top_conv[0][0]']

```

```
top_activation (Activation)      (None, 8, 8, 1280)    0
['top_bn[0][0]']
```

```
=====
=====
Total params: 5,919,312
Trainable params: 5,858,704
Non-trainable params: 60,608
-----
-----
```

```
[ ]: DROPOUT_RATE = 0.6
      L1_PENALTY = 1e-5
      L2_PENALTY = 1e-5

def create_transferred_model(base_model):
    base_model.trainable = False # Freezing the base model

    inputs = keras.Input(shape=(256, 256, 3),
                           batch_size=BATCH_SIZE)
    scaled = keras.layers.Rescaling(scale=1./255.)(inputs)

    x_base = base_model(scaled, training=False)
    gap_layer = keras.layers.GlobalAveragePooling2D()(x_base)

    dropout_layer = keras.layers.Dropout(DROPOUT_RATE)(gap_layer)

    outputs = keras.layers.Dense(NUM_CLASSES,
                                   activation="softmax",
                                   kernel_regularizer=keras.regularizers.
↳L1L2(l1=L1_PENALTY, l2=L2_PENALTY),
                                   )(dropout_layer)

    return keras.models.Model(inputs, outputs)
```

```
[ ]: transferred_model = create_transferred_model(efficient_net)
      transferred_model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(32, 256, 256, 3)]	0
rescaling (Rescaling)	(32, 256, 256, 3)	0
efficientnetv2-b0 (Functional)	(None, 8, 8, 1280)	5919312

global_average_pooling2d (GlobalAveragePooling2D)	(32, 1280)	0
dropout (Dropout)	(32, 1280)	0
dense (Dense)	(32, 20)	25620

```
=====
Total params: 5,944,932
Trainable params: 25,620
Non-trainable params: 5,919,312
-----
```

1.3 Bez transfer learningu

Model bez transfer learningu wykazuje precyzję na poziomie 3%. Jest to równoważne losowaniu, czyli model bez zastosowania transfer learningu jest nieefektywny.

```
[ ]: without_training = transferred_model
without_training.compile(optimizer="adam", loss="categorical_crossentropy",
    ↳metrics=[keras.metrics.CategoricalAccuracy()])
without_training.evaluate(validation_ds)
```

```
32/32 [=====] - 12s 77ms/step - loss: 3.1183 -
categorical_accuracy: 0.0352
```

```
[ ]: [3.118260383605957, 0.03515625]
```

1.4 Transfer learning

W pierwszych próbach transfer learningu model overfitował. Zastosowanie warstwy `dropout` oraz regularyzacji L2 pozwoliło zniwelować ten problem i uzyskać precyzję na poziomie 72% w porównaniu z wcześniejszą rzędu ~50%.

Przed użyciem warstwy dropoutu

Po użyciu warstwy dropoutu

```
[ ]: train_and_check_model(transferred_model, "Transfer learning model",
    ↳t_ds=train_ds, v_ds=validation_ds, epochs=100, lr_scale=0.2)
```

```
Epoch 1/100
128/128 [=====] - 23s 111ms/step - loss: 2.9732 -
categorical_accuracy: 0.1064 - val_loss: 2.5579 - val_categorical_accuracy:
0.3574
Epoch 2/100
128/128 [=====] - 14s 102ms/step - loss: 2.4726 -
```

categorical_accuracy: 0.2827 - val_loss: 2.1727 - val_categorical_accuracy: 0.5088

Epoch 3/100
 128/128 [=====] - 13s 101ms/step - loss: 2.1316 - categorical_accuracy: 0.4016 - val_loss: 1.9176 - val_categorical_accuracy: 0.5674

Epoch 4/100
 128/128 [=====] - 13s 100ms/step - loss: 1.9076 - categorical_accuracy: 0.4697 - val_loss: 1.7433 - val_categorical_accuracy: 0.5879

Epoch 5/100
 128/128 [=====] - 14s 103ms/step - loss: 1.7407 - categorical_accuracy: 0.5168 - val_loss: 1.6189 - val_categorical_accuracy: 0.6094

Epoch 6/100
 128/128 [=====] - 13s 99ms/step - loss: 1.6277 - categorical_accuracy: 0.5547 - val_loss: 1.5245 - val_categorical_accuracy: 0.6279

Epoch 7/100
 128/128 [=====] - 13s 99ms/step - loss: 1.5363 - categorical_accuracy: 0.5642 - val_loss: 1.4520 - val_categorical_accuracy: 0.6396

Epoch 8/100
 128/128 [=====] - 13s 99ms/step - loss: 1.4542 - categorical_accuracy: 0.5938 - val_loss: 1.3964 - val_categorical_accuracy: 0.6475

Epoch 9/100
 128/128 [=====] - 13s 99ms/step - loss: 1.4046 - categorical_accuracy: 0.6050 - val_loss: 1.3521 - val_categorical_accuracy: 0.6523

Epoch 10/100
 128/128 [=====] - 13s 100ms/step - loss: 1.3574 - categorical_accuracy: 0.6155 - val_loss: 1.3122 - val_categorical_accuracy: 0.6572

Epoch 11/100
 128/128 [=====] - 13s 100ms/step - loss: 1.2957 - categorical_accuracy: 0.6387 - val_loss: 1.2806 - val_categorical_accuracy: 0.6572

Epoch 12/100
 128/128 [=====] - 14s 103ms/step - loss: 1.2672 - categorical_accuracy: 0.6350 - val_loss: 1.2525 - val_categorical_accuracy: 0.6592

Epoch 13/100
 128/128 [=====] - 14s 103ms/step - loss: 1.2381 - categorical_accuracy: 0.6479 - val_loss: 1.2303 - val_categorical_accuracy: 0.6602

Epoch 14/100
 128/128 [=====] - 13s 101ms/step - loss: 1.2088 -

categorical_accuracy: 0.6575 - val_loss: 1.2091 - val_categorical_accuracy: 0.6670

Epoch 15/100
 128/128 [=====] - 13s 101ms/step - loss: 1.1768 - categorical_accuracy: 0.6707 - val_loss: 1.1899 - val_categorical_accuracy: 0.6680

Epoch 16/100
 128/128 [=====] - 13s 100ms/step - loss: 1.1601 - categorical_accuracy: 0.6699 - val_loss: 1.1723 - val_categorical_accuracy: 0.6709

Epoch 17/100
 128/128 [=====] - 14s 102ms/step - loss: 1.1386 - categorical_accuracy: 0.6738 - val_loss: 1.1587 - val_categorical_accuracy: 0.6738

Epoch 18/100
 128/128 [=====] - 14s 101ms/step - loss: 1.1116 - categorical_accuracy: 0.6841 - val_loss: 1.1446 - val_categorical_accuracy: 0.6836

Epoch 19/100
 128/128 [=====] - 13s 100ms/step - loss: 1.0906 - categorical_accuracy: 0.6882 - val_loss: 1.1323 - val_categorical_accuracy: 0.6826

Epoch 20/100
 128/128 [=====] - 13s 98ms/step - loss: 1.0675 - categorical_accuracy: 0.6995 - val_loss: 1.1221 - val_categorical_accuracy: 0.6865

Epoch 21/100
 128/128 [=====] - 13s 100ms/step - loss: 1.0584 - categorical_accuracy: 0.6970 - val_loss: 1.1106 - val_categorical_accuracy: 0.6855

Epoch 22/100
 128/128 [=====] - 13s 100ms/step - loss: 1.0521 - categorical_accuracy: 0.6960 - val_loss: 1.1022 - val_categorical_accuracy: 0.6875

Epoch 23/100
 128/128 [=====] - 13s 100ms/step - loss: 1.0329 - categorical_accuracy: 0.6968 - val_loss: 1.0932 - val_categorical_accuracy: 0.6924

Epoch 24/100
 128/128 [=====] - 13s 99ms/step - loss: 1.0251 - categorical_accuracy: 0.7092 - val_loss: 1.0857 - val_categorical_accuracy: 0.6914

Epoch 25/100
 128/128 [=====] - 13s 101ms/step - loss: 0.9958 - categorical_accuracy: 0.7156 - val_loss: 1.0787 - val_categorical_accuracy: 0.6934

Epoch 26/100
 128/128 [=====] - 13s 100ms/step - loss: 0.9997 -

categorical_accuracy: 0.7144 - val_loss: 1.0705 - val_categorical_accuracy: 0.6963

Epoch 27/100
 128/128 [=====] - 13s 99ms/step - loss: 0.9823 - categorical_accuracy: 0.7163 - val_loss: 1.0646 - val_categorical_accuracy: 0.6992

Epoch 28/100
 128/128 [=====] - 13s 101ms/step - loss: 0.9645 - categorical_accuracy: 0.7244 - val_loss: 1.0586 - val_categorical_accuracy: 0.7002

Epoch 29/100
 128/128 [=====] - 14s 102ms/step - loss: 0.9577 - categorical_accuracy: 0.7207 - val_loss: 1.0540 - val_categorical_accuracy: 0.7021

Epoch 30/100
 128/128 [=====] - 13s 100ms/step - loss: 0.9486 - categorical_accuracy: 0.7251 - val_loss: 1.0476 - val_categorical_accuracy: 0.7051

Epoch 31/100
 128/128 [=====] - 13s 100ms/step - loss: 0.9496 - categorical_accuracy: 0.7244 - val_loss: 1.0432 - val_categorical_accuracy: 0.7021

Epoch 32/100
 128/128 [=====] - 14s 102ms/step - loss: 0.9291 - categorical_accuracy: 0.7305 - val_loss: 1.0379 - val_categorical_accuracy: 0.7041

Epoch 33/100
 128/128 [=====] - 13s 100ms/step - loss: 0.9251 - categorical_accuracy: 0.7358 - val_loss: 1.0328 - val_categorical_accuracy: 0.7061

Epoch 34/100
 128/128 [=====] - 13s 99ms/step - loss: 0.9203 - categorical_accuracy: 0.7375 - val_loss: 1.0306 - val_categorical_accuracy: 0.7080

Epoch 35/100
 128/128 [=====] - 14s 102ms/step - loss: 0.9304 - categorical_accuracy: 0.7280 - val_loss: 1.0271 - val_categorical_accuracy: 0.7041

Epoch 36/100
 128/128 [=====] - 13s 99ms/step - loss: 0.9243 - categorical_accuracy: 0.7307 - val_loss: 1.0224 - val_categorical_accuracy: 0.7080

Epoch 37/100
 128/128 [=====] - 13s 100ms/step - loss: 0.9010 - categorical_accuracy: 0.7432 - val_loss: 1.0179 - val_categorical_accuracy: 0.7100

Epoch 38/100
 128/128 [=====] - 14s 103ms/step - loss: 0.8840 -

categorical_accuracy: 0.7529 - val_loss: 1.0155 - val_categorical_accuracy: 0.7070

Epoch 39/100
 128/128 [=====] - 13s 101ms/step - loss: 0.8909 - categorical_accuracy: 0.7402 - val_loss: 1.0122 - val_categorical_accuracy: 0.7090

Epoch 40/100
 128/128 [=====] - 14s 102ms/step - loss: 0.8772 - categorical_accuracy: 0.7407 - val_loss: 1.0092 - val_categorical_accuracy: 0.7100

Epoch 41/100
 128/128 [=====] - 13s 101ms/step - loss: 0.8779 - categorical_accuracy: 0.7454 - val_loss: 1.0073 - val_categorical_accuracy: 0.7100

Epoch 42/100
 128/128 [=====] - 13s 100ms/step - loss: 0.8661 - categorical_accuracy: 0.7524 - val_loss: 1.0050 - val_categorical_accuracy: 0.7148

Epoch 43/100
 128/128 [=====] - 16s 120ms/step - loss: 0.8630 - categorical_accuracy: 0.7542 - val_loss: 1.0029 - val_categorical_accuracy: 0.7139

Epoch 44/100
 128/128 [=====] - 13s 100ms/step - loss: 0.8588 - categorical_accuracy: 0.7476 - val_loss: 0.9977 - val_categorical_accuracy: 0.7129

Epoch 45/100
 128/128 [=====] - 13s 100ms/step - loss: 0.8525 - categorical_accuracy: 0.7542 - val_loss: 0.9980 - val_categorical_accuracy: 0.7148

Epoch 46/100
 128/128 [=====] - 13s 99ms/step - loss: 0.8541 - categorical_accuracy: 0.7507 - val_loss: 0.9948 - val_categorical_accuracy: 0.7158

Epoch 47/100
 128/128 [=====] - 13s 101ms/step - loss: 0.8406 - categorical_accuracy: 0.7600 - val_loss: 0.9924 - val_categorical_accuracy: 0.7168

Epoch 48/100
 128/128 [=====] - 14s 100ms/step - loss: 0.8372 - categorical_accuracy: 0.7559 - val_loss: 0.9903 - val_categorical_accuracy: 0.7197

Epoch 49/100
 128/128 [=====] - 13s 100ms/step - loss: 0.8347 - categorical_accuracy: 0.7600 - val_loss: 0.9878 - val_categorical_accuracy: 0.7217

Epoch 50/100
 128/128 [=====] - 13s 100ms/step - loss: 0.8284 -

categorical_accuracy: 0.7622 - val_loss: 0.9862 - val_categorical_accuracy: 0.7207

Epoch 51/100
 128/128 [=====] - 13s 98ms/step - loss: 0.8138 - categorical_accuracy: 0.7656 - val_loss: 0.9850 - val_categorical_accuracy: 0.7197

Epoch 52/100
 128/128 [=====] - 13s 98ms/step - loss: 0.8219 - categorical_accuracy: 0.7639 - val_loss: 0.9840 - val_categorical_accuracy: 0.7188

Epoch 53/100
 128/128 [=====] - 13s 99ms/step - loss: 0.8219 - categorical_accuracy: 0.7620 - val_loss: 0.9819 - val_categorical_accuracy: 0.7227

Epoch 54/100
 128/128 [=====] - 13s 101ms/step - loss: 0.8203 - categorical_accuracy: 0.7600 - val_loss: 0.9814 - val_categorical_accuracy: 0.7139

Epoch 55/100
 128/128 [=====] - 13s 99ms/step - loss: 0.8185 - categorical_accuracy: 0.7561 - val_loss: 0.9799 - val_categorical_accuracy: 0.7197

Epoch 56/100
 128/128 [=====] - 13s 100ms/step - loss: 0.8034 - categorical_accuracy: 0.7612 - val_loss: 0.9801 - val_categorical_accuracy: 0.7256

Epoch 57/100
 128/128 [=====] - 13s 101ms/step - loss: 0.7904 - categorical_accuracy: 0.7654 - val_loss: 0.9772 - val_categorical_accuracy: 0.7246

Epoch 58/100
 128/128 [=====] - 13s 98ms/step - loss: 0.7979 - categorical_accuracy: 0.7703 - val_loss: 0.9751 - val_categorical_accuracy: 0.7236

Epoch 59/100
 128/128 [=====] - 13s 97ms/step - loss: 0.7794 - categorical_accuracy: 0.7786 - val_loss: 0.9761 - val_categorical_accuracy: 0.7246

Epoch 60/100
 128/128 [=====] - 13s 98ms/step - loss: 0.7936 - categorical_accuracy: 0.7686 - val_loss: 0.9737 - val_categorical_accuracy: 0.7246

Epoch 61/100
 128/128 [=====] - 13s 97ms/step - loss: 0.7757 - categorical_accuracy: 0.7700 - val_loss: 0.9732 - val_categorical_accuracy: 0.7256

Epoch 62/100
 128/128 [=====] - 13s 99ms/step - loss: 0.7876 -

categorical_accuracy: 0.7688 - val_loss: 0.9711 - val_categorical_accuracy: 0.7246

Epoch 63/100
 128/128 [=====] - 13s 100ms/step - loss: 0.7815 - categorical_accuracy: 0.7715 - val_loss: 0.9716 - val_categorical_accuracy: 0.7227

Epoch 64/100
 128/128 [=====] - 13s 98ms/step - loss: 0.7760 - categorical_accuracy: 0.7778 - val_loss: 0.9674 - val_categorical_accuracy: 0.7275

Epoch 65/100
 128/128 [=====] - 13s 99ms/step - loss: 0.7822 - categorical_accuracy: 0.7676 - val_loss: 0.9659 - val_categorical_accuracy: 0.7266

Epoch 66/100
 128/128 [=====] - 13s 100ms/step - loss: 0.7684 - categorical_accuracy: 0.7778 - val_loss: 0.9645 - val_categorical_accuracy: 0.7275

Epoch 67/100
 128/128 [=====] - 13s 98ms/step - loss: 0.7712 - categorical_accuracy: 0.7705 - val_loss: 0.9643 - val_categorical_accuracy: 0.7246

Epoch 68/100
 128/128 [=====] - 15s 115ms/step - loss: 0.7661 - categorical_accuracy: 0.7788 - val_loss: 0.9619 - val_categorical_accuracy: 0.7266

Epoch 69/100
 128/128 [=====] - 13s 98ms/step - loss: 0.7672 - categorical_accuracy: 0.7800 - val_loss: 0.9621 - val_categorical_accuracy: 0.7266

Epoch 70/100
 128/128 [=====] - 13s 98ms/step - loss: 0.7582 - categorical_accuracy: 0.7788 - val_loss: 0.9642 - val_categorical_accuracy: 0.7236

Epoch 71/100
 128/128 [=====] - 13s 99ms/step - loss: 0.7572 - categorical_accuracy: 0.7778 - val_loss: 0.9639 - val_categorical_accuracy: 0.7266

Epoch 72/100
 128/128 [=====] - 13s 99ms/step - loss: 0.7564 - categorical_accuracy: 0.7810 - val_loss: 0.9618 - val_categorical_accuracy: 0.7266

Epoch 73/100
 128/128 [=====] - 13s 100ms/step - loss: 0.7471 - categorical_accuracy: 0.7773 - val_loss: 0.9592 - val_categorical_accuracy: 0.7275

Epoch 74/100
 128/128 [=====] - 13s 99ms/step - loss: 0.7471 -

categorical_accuracy: 0.7812 - val_loss: 0.9598 - val_categorical_accuracy: 0.7275

Epoch 75/100
 128/128 [=====] - 13s 102ms/step - loss: 0.7396 - categorical_accuracy: 0.7786 - val_loss: 0.9576 - val_categorical_accuracy: 0.7275

Epoch 76/100
 128/128 [=====] - 13s 100ms/step - loss: 0.7381 - categorical_accuracy: 0.7817 - val_loss: 0.9563 - val_categorical_accuracy: 0.7256

Epoch 77/100
 128/128 [=====] - 13s 101ms/step - loss: 0.7300 - categorical_accuracy: 0.7832 - val_loss: 0.9569 - val_categorical_accuracy: 0.7275

Epoch 78/100
 128/128 [=====] - 13s 99ms/step - loss: 0.7396 - categorical_accuracy: 0.7871 - val_loss: 0.9554 - val_categorical_accuracy: 0.7266

Epoch 79/100
 128/128 [=====] - 13s 98ms/step - loss: 0.7401 - categorical_accuracy: 0.7908 - val_loss: 0.9561 - val_categorical_accuracy: 0.7275

Epoch 80/100
 128/128 [=====] - 13s 98ms/step - loss: 0.7389 - categorical_accuracy: 0.7793 - val_loss: 0.9551 - val_categorical_accuracy: 0.7246

Epoch 81/100
 128/128 [=====] - 13s 98ms/step - loss: 0.7228 - categorical_accuracy: 0.7852 - val_loss: 0.9541 - val_categorical_accuracy: 0.7266

Epoch 82/100
 128/128 [=====] - 14s 102ms/step - loss: 0.7434 - categorical_accuracy: 0.7781 - val_loss: 0.9529 - val_categorical_accuracy: 0.7295

Epoch 83/100
 128/128 [=====] - 13s 101ms/step - loss: 0.7067 - categorical_accuracy: 0.7969 - val_loss: 0.9524 - val_categorical_accuracy: 0.7266

Epoch 84/100
 128/128 [=====] - 13s 101ms/step - loss: 0.7420 - categorical_accuracy: 0.7788 - val_loss: 0.9516 - val_categorical_accuracy: 0.7295

Epoch 85/100
 128/128 [=====] - 13s 101ms/step - loss: 0.7274 - categorical_accuracy: 0.7925 - val_loss: 0.9514 - val_categorical_accuracy: 0.7275

Epoch 86/100
 128/128 [=====] - 13s 99ms/step - loss: 0.7131 -

categorical_accuracy: 0.7900 - val_loss: 0.9501 - val_categorical_accuracy: 0.7256

Epoch 87/100
128/128 [=====] - 13s 99ms/step - loss: 0.7230 - categorical_accuracy: 0.7947 - val_loss: 0.9514 - val_categorical_accuracy: 0.7266

Epoch 88/100
128/128 [=====] - 13s 98ms/step - loss: 0.7180 - categorical_accuracy: 0.7959 - val_loss: 0.9500 - val_categorical_accuracy: 0.7295

Epoch 89/100
128/128 [=====] - 13s 101ms/step - loss: 0.7124 - categorical_accuracy: 0.7913 - val_loss: 0.9498 - val_categorical_accuracy: 0.7285

Epoch 90/100
128/128 [=====] - 13s 101ms/step - loss: 0.7090 - categorical_accuracy: 0.7896 - val_loss: 0.9512 - val_categorical_accuracy: 0.7305

Epoch 91/100
128/128 [=====] - 14s 104ms/step - loss: 0.7062 - categorical_accuracy: 0.7947 - val_loss: 0.9490 - val_categorical_accuracy: 0.7236

Epoch 92/100
128/128 [=====] - 13s 101ms/step - loss: 0.7165 - categorical_accuracy: 0.7866 - val_loss: 0.9494 - val_categorical_accuracy: 0.7256

Epoch 93/100
128/128 [=====] - 14s 102ms/step - loss: 0.7062 - categorical_accuracy: 0.7913 - val_loss: 0.9485 - val_categorical_accuracy: 0.7246

Epoch 94/100
128/128 [=====] - 13s 102ms/step - loss: 0.7064 - categorical_accuracy: 0.7898 - val_loss: 0.9475 - val_categorical_accuracy: 0.7246

Epoch 95/100
128/128 [=====] - 13s 100ms/step - loss: 0.7024 - categorical_accuracy: 0.7969 - val_loss: 0.9480 - val_categorical_accuracy: 0.7256

Epoch 96/100
128/128 [=====] - 13s 99ms/step - loss: 0.7069 - categorical_accuracy: 0.7930 - val_loss: 0.9446 - val_categorical_accuracy: 0.7246

Epoch 97/100
128/128 [=====] - 13s 99ms/step - loss: 0.7066 - categorical_accuracy: 0.7896 - val_loss: 0.9448 - val_categorical_accuracy: 0.7246

Epoch 98/100
128/128 [=====] - 13s 100ms/step - loss: 0.6978 -

```
categorical_accuracy: 0.7939 - val_loss: 0.9453 - val_categorical_accuracy: 0.7227
```


```
Epoch 99/100
```

```
128/128 [=====] - 14s 102ms/step - loss: 0.7041 - categorical_accuracy: 0.7922 - val_loss: 0.9439 - val_categorical_accuracy: 0.7295
```

```
Epoch 100/100
```

```
128/128 [=====] - 14s 102ms/step - loss: 0.7150 - categorical_accuracy: 0.7854 - val_loss: 0.9444 - val_categorical_accuracy: 0.7266
```

```
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing 5 of 91). These functions will not be directly callable after loading.
```



```
[ ]: transferred_model.evaluate(validation_ds)
```

```
32/32 [=====] - 2s 75ms/step - loss: 0.9444 - categorical_accuracy: 0.7266
```

```
[ ]: [0.944395899772644, 0.7265625]
```

1.4.1 Tablica pomyłek


```
[ ]: predictions = np.array([])
labels = np.array([])
for x, y in validation_ds:
    predictions = np.concatenate([predictions, np.argmax(transferred_model.
↪predict(x), axis=-1)])
    labels = np.concatenate([labels, np.argmax(y, axis=-1)])

cf_matrix = tf.math.confusion_matrix(labels=labels, predictions=predictions).
↪numpy()
sns.heatmap(cf_matrix,
            annot=True,
            cmap='Blues')
```

```
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 52ms/step
1/1 [=====] - 0s 70ms/step
1/1 [=====] - 0s 79ms/step
1/1 [=====] - 0s 57ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 58ms/step
1/1 [=====] - 0s 49ms/step
1/1 [=====] - 0s 48ms/step
1/1 [=====] - 0s 58ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 49ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 52ms/step
1/1 [=====] - 0s 56ms/step
1/1 [=====] - 0s 67ms/step
1/1 [=====] - 0s 60ms/step
1/1 [=====] - 0s 56ms/step
1/1 [=====] - 0s 55ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 62ms/step
1/1 [=====] - 0s 49ms/step
1/1 [=====] - 0s 53ms/step
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 42ms/step
```

1/1 [=====] - 0s 38ms/step

[]: <Axes: >



Powyższa tabela pokazuje, że nasz model dosyć dobrze radzi sobie z zadaniem klasyfikacji. Najwięcej (dokładnie 8) błędów popełnia dla klas nr 6 i 7, czyli *ice cream* i *fish and chips*. Jest to jednak stosunkowo mała liczba w porównaniu z klasyfikacją poprawnych klas, która waha się w między 35 a ponad 50.

1.5 Fine tuning

Następnie zastosowaliśmy fine-tuning, który pozwolił na zwiększenie precyzji do 77% (różnica 5 pp.).

```
[ ]: def fine_tune(model_after_transfer_learning):
    base_model_idx = 2
    base_model = model_after_transfer_learning.get_layer(index=base_model_idx)
    base_model.trainable = True

    train_and_check_model(model_after_transfer_learning,
                           "Fine tuned ConvNeXt",
                           t_ds=train_ds,
                           v_ds=validation_ds,
```

```
epochs=10, lr_scale=0.01)
fine_tune(transferred_model)
```

```
Epoch 1/10
128/128 [=====] - 131s 311ms/step - loss: 0.6556 -
categorical_accuracy: 0.8018 - val_loss: 0.8422 - val_categorical_accuracy:
0.7617
```

```
Epoch 2/10
128/128 [=====] - 37s 287ms/step - loss: 0.5766 -
categorical_accuracy: 0.8347 - val_loss: 0.8179 - val_categorical_accuracy:
0.7656
```

```
Epoch 3/10
128/128 [=====] - 37s 286ms/step - loss: 0.5398 -
categorical_accuracy: 0.8396 - val_loss: 0.7928 - val_categorical_accuracy:
0.7715
```

```
Epoch 4/10
128/128 [=====] - 37s 285ms/step - loss: 0.5003 -
categorical_accuracy: 0.8499 - val_loss: 0.7888 - val_categorical_accuracy:
0.7725
```

```
Epoch 5/10
128/128 [=====] - 37s 286ms/step - loss: 0.4679 -
categorical_accuracy: 0.8640 - val_loss: 0.7819 - val_categorical_accuracy:
0.7734
```

```
Epoch 6/10
128/128 [=====] - 37s 286ms/step - loss: 0.4522 -
categorical_accuracy: 0.8735 - val_loss: 0.7674 - val_categorical_accuracy:
0.7773
```

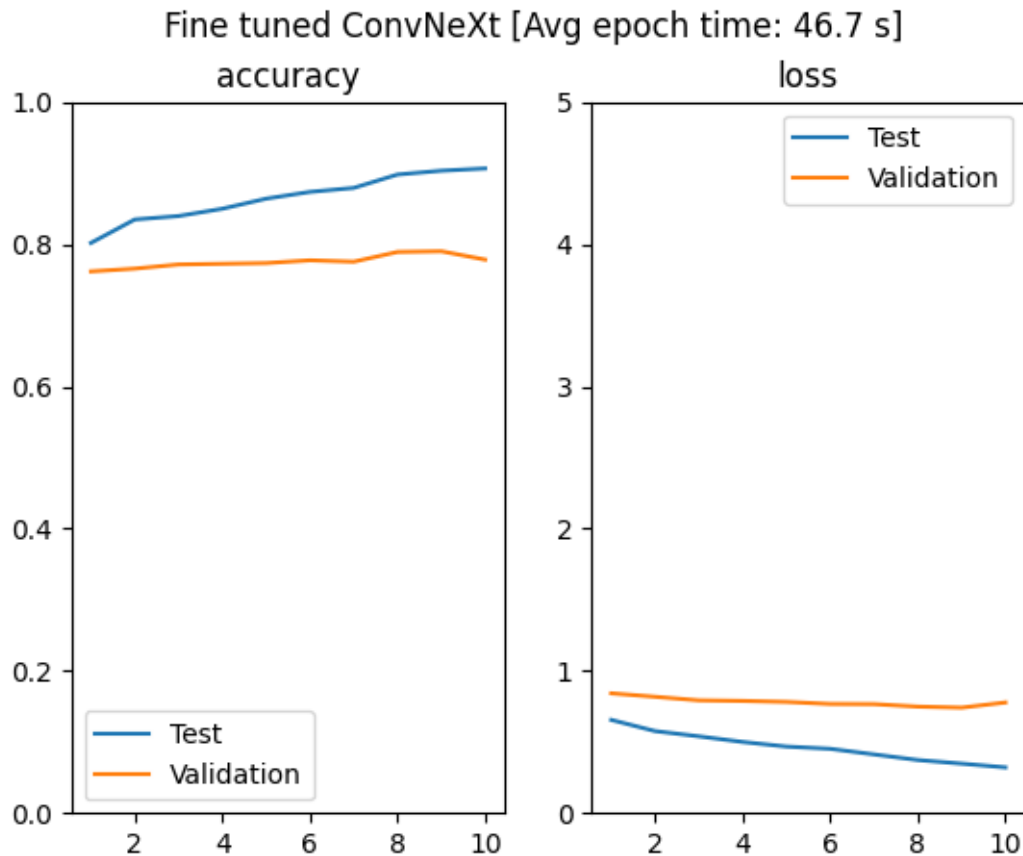
```
Epoch 7/10
128/128 [=====] - 37s 284ms/step - loss: 0.4128 -
categorical_accuracy: 0.8792 - val_loss: 0.7661 - val_categorical_accuracy:
0.7754
```

```
Epoch 8/10
128/128 [=====] - 38s 292ms/step - loss: 0.3731 -
categorical_accuracy: 0.8979 - val_loss: 0.7487 - val_categorical_accuracy:
0.7891
```

```
Epoch 9/10
128/128 [=====] - 38s 292ms/step - loss: 0.3479 -
categorical_accuracy: 0.9033 - val_loss: 0.7416 - val_categorical_accuracy:
0.7900
```

```
Epoch 10/10
128/128 [=====] - 38s 292ms/step - loss: 0.3214 -
categorical_accuracy: 0.9065 - val_loss: 0.7778 - val_categorical_accuracy:
0.7783
```

```
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op,
_jit_compiled_convolution_op while saving (showing
5 of 91). These functions will not be directly callable after loading.
```

```
[ ]: transferred_model.evaluate(validation_ds)

predictions = np.array([])
labels = np.array([])
for x, y in validation_ds:
    predictions = np.concatenate([predictions, np.argmax(transferred_model.
        ↪predict(x), axis=-1)])
    labels = np.concatenate([labels, np.argmax(y, axis=-1)])

cf_matrix = tf.math.confusion_matrix(labels=labels, predictions=predictions).
    ↪numpy()
sns.heatmap(cf_matrix,
            annot=True,
            cmap='Blues')
```


```
32/32 [=====] - 3s 81ms/step - loss: 0.7778 -
categorical_accuracy: 0.7783
1/1 [=====] - 2s 2s/step
1/1 [=====] - 0s 66ms/step
1/1 [=====] - 0s 88ms/step
```

```

1/1 [=====] - 0s 70ms/step
1/1 [=====] - 0s 52ms/step
1/1 [=====] - 0s 57ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 62ms/step
1/1 [=====] - 0s 54ms/step
1/1 [=====] - 0s 56ms/step
1/1 [=====] - 0s 43ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 58ms/step
1/1 [=====] - 0s 61ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 54ms/step
1/1 [=====] - 0s 82ms/step
1/1 [=====] - 0s 103ms/step
1/1 [=====] - 0s 67ms/step
1/1 [=====] - 0s 80ms/step
1/1 [=====] - 0s 77ms/step
1/1 [=====] - 0s 146ms/step
1/1 [=====] - 0s 71ms/step
1/1 [=====] - 0s 96ms/step
1/1 [=====] - 0s 54ms/step
1/1 [=====] - 0s 65ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 39ms/step

```

[]: <Axes: >



Jak należało się spodziewać liczba poprawnych wyników zwiększyła się. Wspomniana wcześniej liczba błędów dla problemu klasyfikacji klas nr 6 i 7 zmalała z 8 do 5. Model znacznie lepiej radzi sobie z trudniejszymi przykładami niż model, na którym nie zastosowano fine-tuningu.

1.6 Wnioski

Transfer learning pozwala na szybkie uczenie, co pozwala zaoszczędzić czas i zasoby. Dzieje się tak, ponieważ wykorzystuje wstępnie wytrenowane modele, które nauczyły się reprezentacji danych na dużych zbiorach danych. Przygotowany w ten sposób model radzi sobie zdecydowanie lepiej niż wytrenowany od zera. Jest to szczególnie istotne w przypadku ograniczonej ilości dostępnych danych. Wyniki powyżej wytrenowanych modeli potwierdzają te wnioski.

Z kolei fine-tuning pozwala na aktualizację parametrów modelu na nowych danych, co umożliwia lepszą adaptację modelu do specyficznych warunków i zmian w danych wejściowych. W niniejszej pracy zastosowanie tej techniki pozwoliło znacząco polepszyć wyniki modelu.